

ECONOMIC COMMISSION FOR LATIN AMERICA AND THE CARIBBEAN

Foundations of Generative AI

Fernando Fetis Riquelme

Spring 2026



NACIONES UNIDAS

CEPAL

LECTURE 1


What this talk covers

- 1 Learning from data
- 2 Language models
- 3 From model to assistant
- 4 Demonstrations
- 5 Beyond text
- 6 Current limitations
- 7 Closing

01

Learning from data


Rule-based and learned systems



KEY IDEA

Experts cannot state what they know. A radiologist sees a tumor in a second and cannot write down the rule.


How to optimize



SGD: it starts at a random point and is iteratively improved until a local minimum is reached.

$$x_{t+1} = x_t - \eta \cdot f'(x_t)$$

Fitting a model to data




$$\text{model}(x) = a_0 + a_1x + a_2x^2 + a_3x^3$$

$$\min_{a_0, \dots, a_3} \sum_n (y_n - \text{model}(x_n))^2$$

The coefficients $a_0, \dots, a_3$ are the **parameters**, which are adjusted in each iteration of SGD to reduce the fitting error.

Neural networks

Modern neural networks are made up of hundreds of simpler building blocks



Each block contains a set of parameters that are adjusted during training

Image classification



dog



cat





dog

... repeated millions of times



one training photo

true label: cat




Still 28% incorrect. This is used to further improve the model.

KEY IDEA

With enough examples a neural network can **learn** essentially any relationship. Nobody programs the rule for “cat”, it falls out of the data.

(Cybenko, 1989)

Position within artificial intelligence



examples:

robotics · planning · expert systems

examples:

spam filters · credit scoring


examples:

face recognition · speech-to-text

examples:

chat assistants · image synthesis

Milestones, 2012–2026



Everything on this line is the recipe from the previous slides, scaled up and pointed at new data.

(Brown et al., 2020; DeepSeek-AI, 2025; Goodfellow et al., 2014; Krizhevsky et al., 2012; Sutskever et al., 2014; Vaswani et al., 2017)

02

Language models


Tokenization

A language model starts by dividing the text into subunits called **tokens**. Then, a neural network processes the sequence of tokens.



An LLM can read a maximum number of tokens within its **context window**. Modern models can read up to 1 million tokens.

Autoregressive generation




The model outputs $\text{Prob}(\text{next token} \mid \text{context})$ for every possible next token, and one is sampled. In particular, the same prompt can give different continuations.

Few-shot prompting


With enough data, autoregressive training learns to repeat certain patterns:

the prompt: two solved examples, one unfinished



This was first observed in GPT 2 and used systematically in GPT 3.

Emergent abilities



THREE-DIGIT ADDITION

Q: 128 + 367 = 495

Q: 512 + 287 = 799

Q: 234 + 567 = 801

WORD UNSCRAMBLE

Q: elppa = apple

Q: dloihya = holiday

Figure 2 of Wei et al. (2022). Whether the jumps are real or an artifact of grading is still argued.

Scaling laws




Figure 1 of Kaplan et al. (2020): the error falls predictably with compute, data and size. That is what justified spending billions.

THE BITTER LESSON

“General methods that leverage computation are ultimately the most effective”. The architecture used today has not changed much from the original GPT, and the most notable advances are mainly due to the increase in scale.

Chain-of-thought reasoning

DIRECT ANSWER

A juggler has 16 balls. Half are golf balls, and half of the golf balls are blue. How many blue golf balls are there?


A: 8 ✗

THINKING FIRST

A: $16 / 2 = 8$ golf balls.

Half of them are blue, so

$8 / 2 = 4$. The answer is 4 ✓




DeepSeek-R1: rewarded only on final answers, the model learns to think first, and competition-math accuracy climbs.

Since 2024 models are trained to reason by default. On graduate-level science questions (GPQA) they now match PhD experts in the field.

03

From model to assistant

The three training stages



(Brown et al., 2020; Christiano et al., 2017; Ouyang et al., 2022; Radford et al., 2019)

Tool use

the user asks

What is the population of Chile?

the model writes

Let me check the latest census.

<tool>

search("Chile population")

</tool>

↑
generation
pauses here

the program runs it

runs the search and pastes the output back into the text

the model continues


<result>

19.6 million (2024)

</result>

Chile has about 19.6 million people.

Agents



(Anthropic, 2024; Yao et al., 2023)

04

Demonstrations

Demonstration 1: a chat assistant

THE DATA

12 files dragged into the chat.

Excel and CSV mixed, 21 countries,
2000 to 2025.

Nothing lines up.

What we ask for

1. Audit them: what would break a merge?
2. One clean table.
3. Poverty against inequality.
4. A dashboard to explore it.
5. Where does each number come from?

Demonstration 2: an agent

THE SAME DATA

Nothing is uploaded. The agent opens the folder itself.


What we ask for

1. The same audit, and a table that rebuilds itself.
2. A website: map, a page per country, comparison.
3. A chat over the data, driven by a language model.
4. The findings, and what **cannot** be concluded.
5. A short slide deck, compiled.
6. Test it before calling it done.

05

Beyond text

Generative adversarial networks




every time a fake is caught, the generator improves, until the fakes pass



Faces synthesized by a GAN after training. None of these people exist.

(Goodfellow et al., 2014; Karras et al., 2018)

Diffusion models



Diffusion models constitute the **state of the art** for images, alongside flow matching and autoregressive generation.

Image editing



Inpainting: regenerate only a masked region.



Style transfer: the same scene, re-rendered.

The same machinery allows for coloring photographs, increasing resolution, extending an image beyond its edges, and interpolating between two images.

Audio generation

To generate audio, a similar approach to autoregressive text generation can be used. Alternatively, the GAN or diffusion approach could be employed.



the same loop as text: a short slice of sound becomes a piece, and the model predicts the next one

Video generation

Diffusion models for images can be extended to videos, where the main difference is that a time variable must be added to the model.

Source Video



Edit: Add tinsel streamers to the lantern bottom



Editing a video from a text instruction. Improving fast; still weak on physics and on text inside the frame.

World models

The video generation approach can be extended to **world models** that respond, frame by frame, to actions performed by the user.



DOOM with no game engine: each frame is generated from the previous frames and the player's action.

Robotics

The planning mechanism and tool usage of LLM-based agents can be extended to robotics, where actions can be executed in a physical world.



Step by Step Instruction

1. Move the hand over the red cup
2. Grasp the bottle
3. Tilt the hand to pour the contents from the red cup
4. Return the red cup to the upright position and release it

Action Plan

```
move_hand(align with the bottle)
grasp_object(bottle)
move_hand(above the cup)
release_object(bottle)
```

Target Object: red cup

Environment State

<red cup> <on the table, upright, presumably with contents>

↓


<red cup> <on the table, empty>

Matching Score: 8/10

A language model turns “pour” into a step-by-step action plan the robot executes, then checks the result.

Scientific applications

The same approaches can be used to predict weather or other geospatial phenomena.




GraphCast produces a 10-day global forecast in under a minute, more accurately than the leading conventional system (Science, 2023).

Related models now prove theorems, solve partial differential equations, and flag anomalies in astronomical surveys.

Molecules and materials

Generative models can also be applied to the generation of molecules, where a model can learn to generate new drugs or materials with certain properties.



The diffusion recipe from images, applied to molecules: from noise to a candidate drug.

KEY IDEA

Protein misfolding underlies many diseases, and determining one structure experimentally can take months of laboratory work. **AlphaFold** predicts one in minutes (2024 Nobel Prize in Chemistry).

06

Current limitations

Attention over long contexts




Figure 1 of Liu et al. (2023).

LLMs have trouble finding information in intermediate sections of the context. Mitigating problems of this kind is an active area of research.


IN PRACTICE

Place specific and important instructions at the beginning or end of the prompt. This tip may become obsolete in the near future.

Hallucination, bias and sycophancy



Midjourney, 2023: asked for a white man robbing a store.



Gemini, 2024: asked for a 1943 German soldier.

Hallucination. Plausible and reliable content, but incorrect.

Bias. Gender, race, religion, age, disability, politics.


Sycophancy. Models tend to say what the user wants to hear, and to agree with them even if they are wrong.

Prompt injection

It is possible to construct malicious prompts to encourage the LLM to do things that are outside of their guidelines.

CAUTION

A model that reads a page or a document also reads instructions hidden inside it. White text in a PDF can be enough.



Attacks of this type can cause an LLM to reveal their **system prompt** or provide information about their training.


07

Closing

Five takeaways

- | | | |
|---|--|-----------|
| 1 | A model is fitted, not programmed : training only reduces an error score on examples. | section 1 |
| 2 | Guessing the next token, at scale , is enough to produce language, reasoning and abilities nobody wrote down. | section 2 |
| 3 | The assistant is a thin layer on top: alignment makes it useful, tools and a loop make it act. | section 3 |
| 4 | The same recipe now generates images, audio, video, forecasts and molecules . | section 5 |
| 5 | Fluency is not accuracy: hallucination, bias and prompt injection make the human check part of the work. | section 6 |

Recommended reading



- Generative AI blog: fernandofetis.ai/blog
- Contact: fernando.fetis@uchile.cl

Thank you

Lecture 1 · Foundations of Generative AI



References

- Alakhdar, A., Poczos, B., & Washburn, N. (2024). Diffusion Models in De Novo Drug Design. *Journal of Chemical Information and Modeling*.
- Anthropic. (2024, November 25). *Introducing the Model Context Protocol*. <https://modelcontextprotocol.io/>
- Brooks, T., Peebles, B., & Holmes, C. (2024). *Video Generation Models as World Simulators*.
- Brown, T. B., Mann, B., & Ryder, N. (2020). Language Models Are Few-Shot Learners. *Advances in Neural Information Processing Systems*.
- Christiano, P. F., Leike, J., Brown, T. B., Martic, M., Legg, S., & Amodei, D. (2017). Deep Reinforcement Learning from Human Preferences. *Advances in Neural Information Processing Systems*.
- Cybenko, G. (1989). Approximation by Superpositions of a Sigmoidal Function. *Mathematics of Control, Signals and Systems*, 2, 303–314.
- David, D. (2023,). *AI Sees White Man, But It's Actually a Black Man*. <https://medium.com/utopian/ai-sees-white-man-but-its-actually-a-black-man-abb91990b1e>
- DeepSeek-AI. (2025). *DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning*.
- Gallegos, I. O., Rossi, R. A., & Barrow, J. (2024). Bias and Fairness in Large Language Models: A Survey. *Computational Linguistics*, 50, 1097–1179.
- Goodfellow, I. J., Pouget-Abadie, J., & Mirza, M. (2014). Generative Adversarial Nets. *Advances in Neural Information Processing Systems*.
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., & Fritz, M. (2023). Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection. *ACM Workshop on Artificial Intelligence and Security*.
- Ha, D., & Schmidhuber, J. (2018). *World Models*.
- Ho, J., Jain, A., & Abbeel, P. (2020). Denoising Diffusion Probabilistic Models. *Advances in Neural Information Processing Systems*.
- Jiang, F., Xu, Z., & Niu, L. (2024). ArtPrompt: ASCII Art-Based Jailbreak Attacks Against Aligned LLMs. *Annual Meeting of the Association for Computational Linguistics*.
- Jumper, J., Evans, R., & Pritzel, A. (2021). Highly Accurate Protein Structure Prediction with AlphaFold. *Nature*, 596, 583–589.
- Kalai, A. T., Nachum, O., Vempala, S. S., & Zhang, E. (2025). *Why Language Models Hallucinate*.
- Kaplan, J., McCandlish, S., & Henighan, T. (2020). *Scaling Laws for Neural Language Models*.
- Karras, T., Aila, T., Laine, S., & Lehtinen, J. (2018). Progressive Growing of GANs for Improved Quality, Stability, and Variation. *International Conference on Learning Representations*.
- Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., & Iwasawa, Y. (2022). Large Language Models are Zero-Shot Reasoners. *Advances in Neural Information Processing Systems*.
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). ImageNet Classification with Deep Convolutional Neural Networks. *Advances in Neural Information Processing Systems*.
- Lam, R., Sanchez-Gonzalez, A., & Willson, M. (2023). Learning Skillful Medium-Range Global Weather Forecasting. *Science*, 382, 1416–1421.
- Lipman, Y., Chen, R. T. Q., Ben-Hamu, H., Nickel, M., & Le, M. (2023). Flow Matching for Generative Modeling. *International Conference on Learning Representations*.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., & Liang, P. (2023). Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics*.

- Nasr, M., Carlini, N., & Hayase, J. (2023). *Scalable Extraction of Training Data from (Production) Language Models*.
- Oord, A. van den, Dieleman, S., & Zen, H. (2016). *WaveNet: A Generative Model for Raw Audio*.
- Ouyang, L., Wu, J., & Jiang, X. (2022). Training Language Models to Follow Instructions with Human Feedback. *Advances in Neural Information Processing Systems*.
- Polyak, A., Zohar, A., & Brown, A. (2024). *Movie Gen: A Cast of Media Foundation Models*.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., & Sutskever, I. (2019). *Language Models Are Unsupervised Multitask Learners*.
- Rein, D., Hou, B. L., & Stickland, A. C. (2023). *GPQA: A Graduate-Level Google-Proof Q&A Benchmark*.
- Robertson, A. (2024, February 21). *Google Apologizes for 'Missing the Mark' After Gemini Generated Racially Diverse Nazis*. <https://www.theverge.com/2024/2/21/24079371/google-ai-gemini-generative-inaccurate-historical>
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2022). High-Resolution Image Synthesis with Latent Diffusion Models. *IEEE Conference on Computer Vision and Pattern Recognition*.
- Schaeffer, R., Miranda, B., & Koyejo, S. (2023). Are Emergent Abilities of Large Language Models a Mirage?. *Advances in Neural Information Processing Systems*.
- Schick, T., Dwivedi-Yu, J., & Dessì, R. (2023). Toolformer: Language Models Can Teach Themselves to Use Tools. *Advances in Neural Information Processing Systems*.
- Sennrich, R., Haddow, B., & Birch, A. (2016). Neural Machine Translation of Rare Words with Subword Units. *Annual Meeting of the Association for Computational Linguistics*.
- Shannon, C. E. (1948). A Mathematical Theory of Communication. *Bell System Technical Journal*, 27(3), 379–423.
- Sharma, M., Tong, M., & Korbak, T. (2023). *Towards Understanding Sycophancy in Language Models*.
- Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., & Poole, B. (2021). Score-Based Generative Modeling through Stochastic Differential Equations. *International Conference on Learning Representations*.
- Sutskever, I., Vinyals, O., & Le, Q. V. (2014). Sequence to Sequence Learning with Neural Networks. *Advances in Neural Information Processing Systems*.
- Sutton, R. S. (2019, March 13). *The Bitter Lesson*. <http://www.incompleteideas.net/IncIdeas/BitterLesson.html>
- Valevski, D., Leviathan, Y., Arar, M., & Fruchter, S. (2024). *Diffusion Models Are Real-Time Game Engines*.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention Is All You Need. *Advances in Neural Information Processing Systems*.
- Wang, J., Wu, Z., & Li, Y. (2024). *Large Language Models for Robotics: Opportunities, Challenges, and Perspectives*.
- Wei, J., Tay, Y., & Bommasani, R. (2022a). Emergent Abilities of Large Language Models. *Transactions on Machine Learning Research*.
- Wei, J., Wang, X., & Schuurmans, D. (2022b). Chain-of-Thought Prompting Elicits Reasoning in Large Language Models. *Advances in Neural Information Processing Systems*.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing Reasoning and Acting in Language Models. *International Conference on Learning Representations*.
- Zhu, J.-Y., Park, T., Isola, P., & Efros, A. A. (2017). Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks. *IEEE International Conference on Computer Vision*.