

held-out perplexity

cheap, runs
during training

says nothing
about usefulness

public benchmarks

comparable
across models

the answers leak
into training data

a model as judge

scales to
thousands of answers

prefers long,
confident answers

people, on your task

the only measure
that decides

slow, and expensive
to repeat