

the model writes its work before answering

generated, and billed

low effort

thinking

wrong

120 tokens

medium effort

thinking

wrong

900 tokens

high effort

thinking

right

4 200 tokens

the same
question

same weights, same prompt: only the budget changes